

# AWS Introduction

## Most Important AWS Topics/Services for Data Science

Here's a curated list of **essential AWS services** specifically for Data Science, with beginner-friendly explanations:

AWS Service	Why It Matters (Simple Explanation)
<b>S3 (Simple Storage Service)</b>	Like a big online USB drive. You use it to store datasets, CSVs, models, logs, etc.
<b>EC2 (Elastic Compute Cloud)</b>	Like renting a remote computer in the cloud where you can run Jupyter, Python code, train models, etc.
<b>SageMaker</b>	AWS's main tool for Data Science. Lets you build, train, deploy ML models without worrying about hardware or backend setup.
<b>IAM (Identity &amp; Access Management)</b>	Like a lock & key system. Controls who can access your AWS stuff (important for security).
<b>Lambda</b>	Run Python code automatically when something happens (like when new data is uploaded). No server needed.
<b>CloudWatch</b>	Like a security camera + diary. It watches your AWS resources, tracks usage, errors, etc.
<b>RDS (Relational Database Service)</b>	Managed SQL databases in the cloud. Think MySQL/PostgreSQL without needing to install anything.
<b>Glue</b>	Helps move and transform data between systems. Useful for cleaning and preparing large datasets.
<b>Athena</b>	Query large CSV/Parquet/JSON files stored in S3 using SQL. No need to move data anywhere.
<b>Step Functions</b>	Like a flowchart engine. Helps automate workflows like "first clean data → then train model → then email results".

AWS Service	Why It Matters (Simple Explanation)
<b>CloudFormation</b>	Like a recipe. You write code to describe your AWS setup, and AWS builds it for you automatically.
<b>ECR + ECS</b> (Optional later)	For containerized apps. Helpful if you use Docker for model deployment. We'll revisit this later.

1. **Amazon S3** – Your cloud "hard drive" for storing datasets, models, and results.  
*Why?* All data science starts with data, and S3 is where you'll store it.
2. **AWS Lambda** – Run code without managing servers (e.g., preprocess data automatically).  
*Why?* Automate small tasks like cleaning data or triggering ML pipelines.
3. **Amazon SageMaker** – Fully managed service to build, train, and deploy ML models.  
*Why?* The heart of AWS ML—no need to set up servers or frameworks manually.
4. **AWS Glue** – Prepares and transforms data for analysis (ETL: Extract, Transform, Load).  
*Why?* Real-world data is messy; Glue cleans it for you.
5. **Amazon EC2** – Virtual computers in the cloud (for heavy computations).  
*Why?* When you need more power than your laptop for training big models.
6. **AWS IAM** – Controls who can access your AWS resources (security first!).  
*Why?* Avoid accidental data leaks or unauthorized costs.
7. **Amazon Athena** – Query data in S3 using SQL (no database setup needed).  
*Why?* Quickly analyze large datasets without moving them.
8. **Amazon QuickSight** – Create dashboards and visualize your data.  
*Why?* Share insights with non-technical teams.
9. **AWS Step Functions** – Coordinate multiple AWS services into workflows.  
*Why?* Chain data preprocessing → training → deployment automatically.

10. **Bedrock (for GenAI)** – Access to foundation models like Claude, Llama 2.

*Why?* Add generative AI (text, images) to your apps without training from scratch.

## Simple Questions

Which topic excites you the most? Or ask:

- "How do I store a CSV file in AWS?" → **S3**
- "How do I train a model without installing anything?" → **SageMaker**
- "How do I automate a data pipeline?" → **Lambda + Glue**

## Common Bill Shock Scenarios (And Fixes)

Scenario	Why It Happens	How to Prevent
EC2 left running 24/7	Forgot to turn off a test server	Use <b>AWS Instance Scheduler</b> or set auto-termination timers
S3 bucket with public access	Accidentally exposed files	Enable <b>S3 Block Public Access</b>
Lambda functions spamming	Infinite loop in code	Set <b>Concurrency Limits</b>

## Key Tips for Beginners

1. **Always enable MFA (Multi-Factor Auth):** Prevents hackers from mining Bitcoin on your account.
2. **Delete unused resources:** Check **AWS Cost Explorer** monthly for "ghost" services.
3. **Use the Free Tier wisely:** Stick to **t2.micro** EC2 instances and small S3 buckets.
4. **Mock bills with the Pricing Calculator:** [AWS Pricing Calculator](#).

## Reality Check

- AWS **won't charge you without warning** (you'll get alerts first).
- Small experiments (e.g., running a server for 1 hour) cost pennies.
- Major leaks usually happen due to **misconfigured auto-scaling** or **public data transfers**.

## What Is the AWS Free Tier?

AWS offers a **12-month Free Tier** for most services.

Examples:

- **S3**: 5 GB of storage free
- **EC2**: 750 hours/month of micro instance (t2.micro or t3.micro)
- **SageMaker**: 250 hours of ml.t2.medium
- **Bedrock**: limited free LLM tokens for some models



But: If you go beyond the free limits (like 6 GB S3 instead of 5 GB), then you're charged for the **extra 1 GB only**.

## Pro Tips for New Users

Mistake	Prevention
Leaving EC2 running 24/7	Stop or terminate it after use
Choosing large instance types	Stick to <code>t2.micro</code> , <code>t3.micro</code> for Free Tier
Not deleting SageMaker notebooks	Stop the notebook to avoid surprise charges
Using services not in Free Tier	Double-check AWS Free Tier website

## What is an EC2 Instance?

Imagine an **EC2 instance** as a "**virtual computer**" that you rent from AWS.

- It's like a **remote laptop** with CPU, RAM, storage, and OS (Linux/Windows).
- You control it entirely: install software, run code, or host websites.
- You pay only for the time it's running (**per second/hour**).
- You can **log in, run code, train ML models, host apps**, etc.

## 💡 Real-Life Example:

🔧 **Case: You want to train a model on 1 million rows of data but your laptop is too slow.**

- You create an EC2 instance with more memory and CPU (or GPU).
- Upload your data and training script.
- Run the training on EC2.
- Download the model weights.
- **Stop the instance.**

Cost? Just the hours you used it — could be ₹10 or ₹100 depending on size/time.

## 🚫 Common Beginner Mistakes


Mistake	Fix
Leaving EC2 running	Stop it after use from AWS console
Picking big instance types	Start with Free Tier ( <code>t2.micro</code> )
Not setting SSH key	You won't be able to login — always download the <code>.pem</code> key safely
Using EC2 to store data permanently	EC2 storage gets erased on termination — store important data in S3

## 🔑 Smart Strategy for Beginners

Use Case	What to Do
Learning EC2, Jupyter, Python setup	✅ Use <code>t2.micro</code> or <code>t3.micro</code> , stop it when done

Use Case	What to Do
Hosting website/chatbot (low traffic)	✓ Use <code>t3.micro</code> , set a <b>₹500 budget alert</b> , monitor usage
Hosting only during the day	✓ Use <b>Lambda or Fargate</b> (event-based, cheaper)
Need GPU for model training	✓ Use EC2 <b>just for the few hours you train</b> , then stop

## Cost Insight

Setup	Monthly Cost (est.)
Single <code>t3.medium</code> instance	₹1,300–₹1,600
Two <code>t3.micro</code> instances	₹1,200–₹1,500 combined
One <code>g4dn.xlarge</code> GPU instance (training only)	₹12,000–₹20,000+ 

# Types of Cloud Infrastructure Services

- Infrastructure as a Service (IaaS)
- Platform as a Service (PaaS)
- Software as a Service (SaaS)

## 1. IaaS (Infrastructure as a Service)



- **Analogy:** You get the land and foundation.

### Your responsibility:

- Operating system (Linux, Windows)
- All your software (Python, libraries, specific versions)
- Applications (Jupyter notebooks, custom data pipelines)
- Network configuration *within* your virtual network
- Security patching *of your OS and software*

### AWS's responsibility:

- Physical servers, networking hardware
- Virtualization layer (making one big server act like many small ones)
- Power, cooling, physical security of data centers

### Key Characteristics:

- **Maximum control:** You decide everything.
- **Flexibility:** Tailor environments exactly to your needs.
- **Cost-effective for variable loads:** Pay only for what you use.
- **Requires expertise:** You need to know how to manage servers.

### AWS Examples for Data Science:

- **Amazon EC2 (Elastic Compute Cloud):** Your virtual machine. Spin up a powerful EC2 instance with many CPUs or GPUs for training complex ML models. Install specific deep learning frameworks (TensorFlow, PyTorch) directly onto it.
- **Amazon S3 (Simple Storage Service):** Object storage for massive datasets. Store raw data lakes, processed data, model artifacts, and project files. Highly durable and scalable, no server to manage for S3 itself.

- **Amazon EBS (Elastic Block Store):** Persistent storage volumes attached to EC2 instances. Like a hard drive for your virtual server, good for databases or when your EC2 needs fast, dedicated storage.
- **Amazon VPC (Virtual Private Cloud):** Your isolated network in the cloud. Define your own IP addresses, subnets, and security rules. Essential for securing your data science environments.
- **Why:** Max control, highly flexible.

## 2. PaaS (Platform as a Service)

- **Analogy:** You get a partially built house.



### Your responsibility:

- Your application code (e.g., Python script for a model API)
- Your data
- Configuration specific to your application

### AWS's responsibility:

- All IaaS components (servers, network, OS)
- Runtime environment (e.g., Python interpreter, Node.js)



- Middleware (web servers like Apache/Nginx, database software)
- Scalability and load balancing for the platform
- Security patching for the platform itself

## Key Characteristics:

- **Faster development/deployment:** Focus on coding, not infrastructure.
- **Managed environment:** AWS handles most operational tasks.
- **Built-in scalability:** Often scales automatically or with minimal config.
- **Less control/flexibility:** Limited access to the underlying infrastructure.


## AWS Examples for Data Science:

- **Amazon RDS (Relational Database Service):** Managed relational databases (PostgreSQL, MySQL, SQL Server, etc.). You don't install or patch the database server; AWS handles it. You just connect and query. Great for structured data.
- **AWS Lambda (Serverless Compute):** Run your code without managing servers. Upload a Python function to process data, trigger it when a new file lands in S3 or on a schedule. Pay only when your code runs. Perfect for event-driven tasks.
- **Amazon SageMaker:** A comprehensive ML platform. Provides managed environments for:
  - **Jupyter Notebooks (SageMaker Studio):** Pre-configured, powerful notebooks for exploration and development.
  - **Training Jobs:** Run large-scale model training without setting up servers.
  - **Model Endpoints:** Deploy models as APIs for real-time predictions; SageMaker manages the hosting.
- **AWS Elastic Beanstalk:** Deploy web applications (like a Flask app serving an ML model) easily. AWS handles the underlying EC2 instances, load balancing, and auto-scaling.

- **Why:** Focus on coding, faster deployment.

### 3. SaaS (Software as a Service)

- **Analogy:** You rent a furnished apartment.



#### Your responsibility:

- Using the application.
- Your data *within* the application.

#### AWS's responsibility:

- Everything: the application itself, its data, runtime, middleware, OS, servers, network, storage.
- All maintenance, updates, security.

#### Key Characteristics:

- **Easiest to use:** Just log in and go.
- **No maintenance burden:** Provider handles everything.
- **Subscription-based:** Usually paid per user or per usage.
- **Least control/customization:** You use it as it's provided.

## AWS Examples (less direct "infrastructure" but good to know):

- **Amazon QuickSight:** Business intelligence service. Connects to your data sources (like S3, RDS) and lets you build interactive dashboards and reports. You don't manage any servers for QuickSight; you just use the web interface.
- **AWS Management Console:** The web interface you use to manage all your AWS services. It's a SaaS application itself!
- **Why:** Easiest to use, no maintenance.

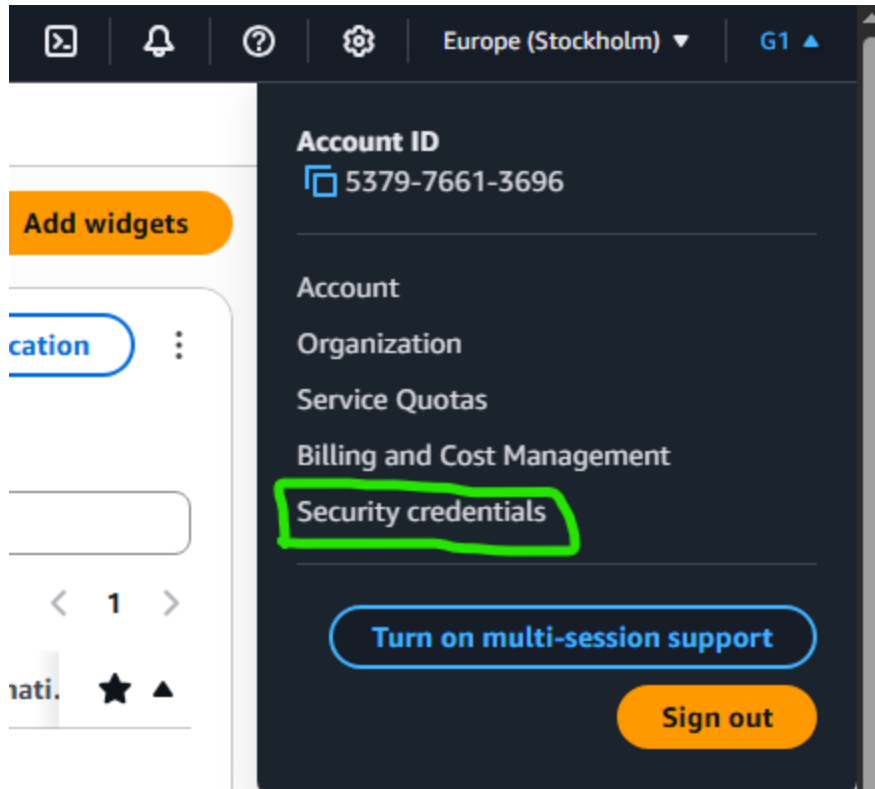
## Bonus: Limit Costs Programmatically

You can even set up a **Lambda function** + **CloudWatch Alarm** to:

- Stop EC2 or SageMaker automatically when you hit budget limits
- Useful if you want to enforce hard spending stops

## MFA (Multi-Factor Authentication) on your AWS Root Account

1. Login console as root user
2. Click your name in upper right hand corner → Security Credentials



### 3. Assign MFA

**My security credentials** [Root user](#) [Info](#)

The root user has access to all AWS resources in this account, and we recommend following [best practices](#). To learn more about the types of AWS credentials and how they're used, see [AWS Security Credentials](#) in AWS General Reference.

**You don't have MFA assigned**  
As a security best practice, we recommend you assign MFA.

[Assign MFA](#)

### 4. Choose Authenticator App

## MFA device

### Device options

In addition to username and password, you will use this device to authenticate into your account.



#### Passkey or security key

Authenticate using your fingerprint, face, or screen lock. Create a passkey on this device or use another device, like a FIDO2 security key.



#### Authenticator app

Authenticate using a code generated by an app installed on your mobile device or computer.



#### Hardware TOTP token

Authenticate using a code generated by Hardware TOTP token or other hardware devices.

## 5. Set up device

- Show QR code
- Scan on Google Auth App

**Set up device** Info

### Authenticator app

A virtual MFA device is an application running on your device that you can configure by scanning a QR code.

- 1

Install a compatible application such as Google Authenticator, Duo Mobile, or Authy app on your mobile device or computer.

[See a list of compatible applications](#)
- 2

Show QR code

Open your authenticator app, choose **Show QR code** on this page, then use the app to scan the code. Alternatively, you can type a secret key. [Show secret key](#)
- 3

Type two consecutive MFA codes below

**Enter a code from your virtual app below**

MFA Code 1

**Wait 30 seconds, and enter a second code entry.**

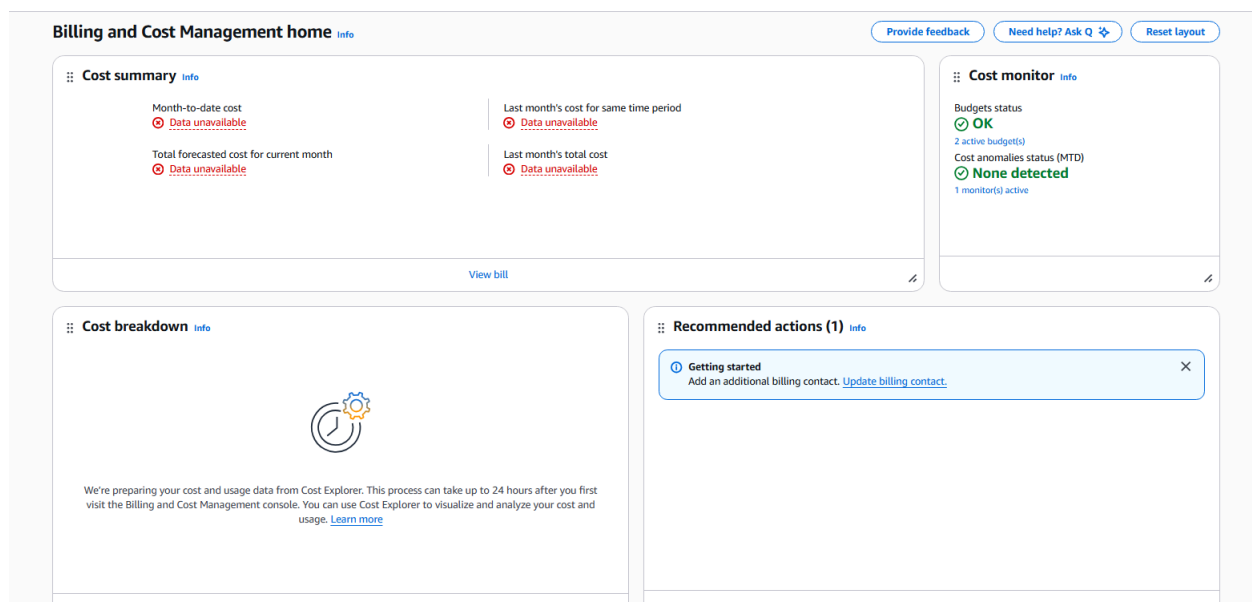
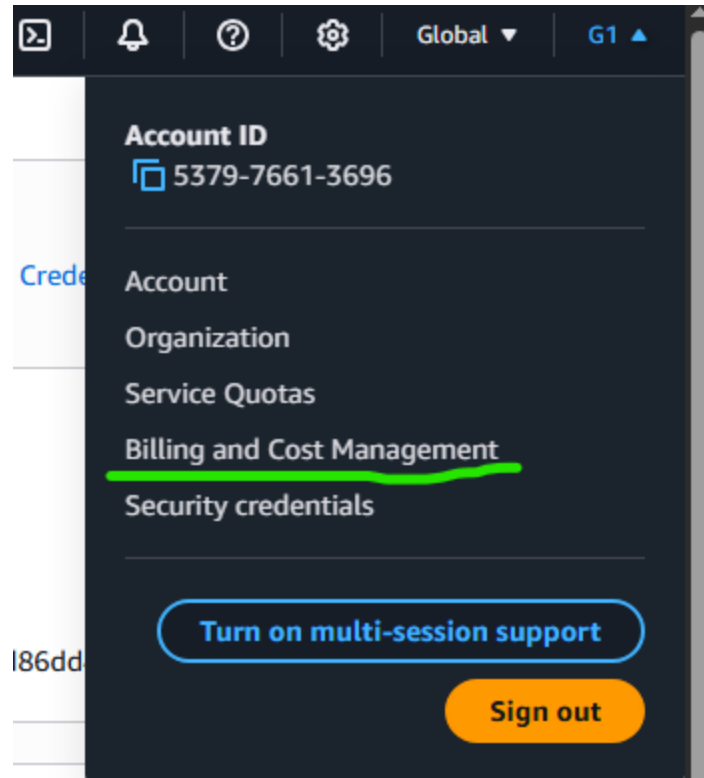
MFA Code 2

[Cancel](#) [Previous](#) [Add MFA](#)

- Enter 2 codes
- Click Add MFA

## Check the Bill

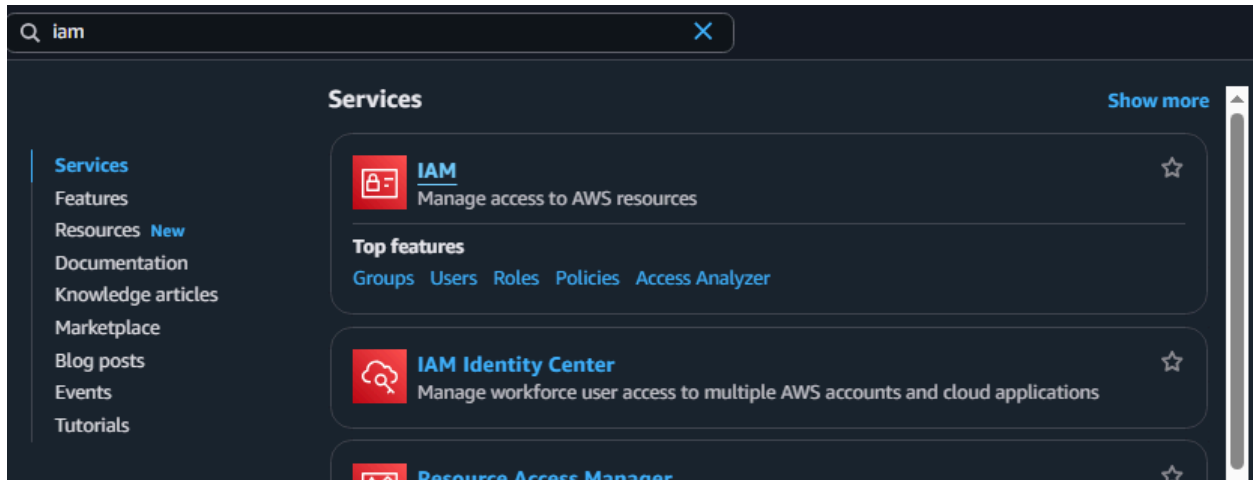
- Login to console (Root)



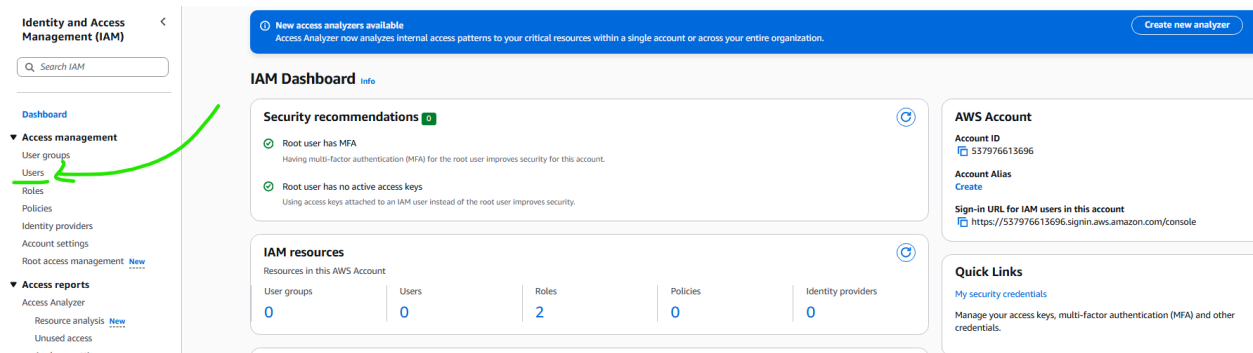
## ► Create an IAM User (!!Extremely Fookin' IMP)

- It's Global service

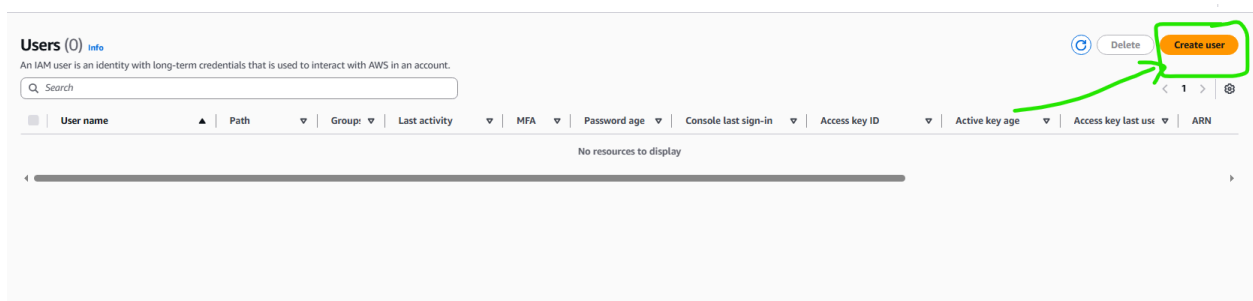
## 1. Go to IAM



## 2. Click Users



## 3. Create User





1. Tick: Provide user access to the AWS Management Console
2. I want to create an IAM user
3. Custom password
  - Enter Password
4. Untick: Users must create a new password at next sign-in - Recommended

**User details**  
**User name**  
  
The user name can have up to 64 characters. Valid characters: A-Z, a-z, 0-9, and + = , . @ \_ - (hyphen)  
☒ **Provide user access to the AWS Management Console - optional**  
If you're providing console access to a person, it's a [best practice](#) to manage their access in IAM Identity Center.

**Are you providing console access to a person?**  
**User type**  
☐ **Specify a user in Identity Center - Recommended**  
We recommend that you use Identity Center to provide console access to a person. With Identity Center, you can centrally manage user access to their AWS accounts and cloud applications.  
☒ **I want to create an IAM user**  
We recommend that you create IAM users only if you need to enable programmatic access through access keys, service-specific credentials for AWS CodeCommit or Amazon Keyspaces, or a backup credential for emergency account access.

**Console password**  
☐ **Autogenerated password**  
You can view the password after you create the user.  
☒ **Custom password**  
Enter a custom password for the user.  
  

- Must be at least 8 characters long
- Must include at least three of the following mix of character types: uppercase letters (A-Z), lowercase letters (a-z), numbers (0-9), and symbols ! @ # \$ % ^ & \* ( ) \_ + - (hyphen) = { } | ' "

☒ **Show password**  
☐ **Users must create a new password at next sign-in - Recommended**  
Users automatically get the [IAMUserChangePassword](#) policy to allow them to change their own password.

! If you are creating programmatic access through access keys or service-specific credentials for AWS CodeCommit or Amazon Keyspaces, you can generate them after you create this IAM user.  
[Learn more](#)

[Cancel](#) [Next](#)

## 4. Set Permissions

## Set permissions

Add user to an existing group or create a new one. Using groups is a best-practice way to manage user's permissions by job functions. [Learn more](#)

### Permissions options

☒ Add user to group

Add user to an existing group, or create a new group. We recommend using groups to manage user permissions by job function.

☐ Copy permissions

Copy all group memberships, attached managed policies, and inline policies from an existing user.

☐ Attach policies directly

Attach a managed policy directly to a user. As a best practice, we recommend attaching policies to a group instead. Then, add the user to the appropriate group.

**1 Get started with groups**

Create a group and select policies to attach to the group. We recommend using groups to manage user permissions by job function, AWS service access, or custom permissions. [Learn more](#)

Create group

### ▼ Set permissions boundary - optional








Set a permissions boundary to control the maximum permissions for this user. Use this advanced feature used to delegate permission management to others. [Learn more](#)

☒ Use a permissions boundary to control the maximum permissions

You can select one of the existing permissions policies to define the boundary.

### Permissions policies (1359)

Select policy to set the permissions boundary.

Filter by Type				
<input type="text" value="Search"/>		All types	< 1 2 3 4 5 6 7 ... 68 > ⚙	
<input type="checkbox"/>	Policy name	Type	Attached entities	
<input type="radio"/>	 <a href="#">AccessAnalyzerServiceRolePolicy</a>	AWS managed	0	
<input type="radio"/>	 <a href="#">AdministratorAccess</a>	AWS managed - job function	0	
<input type="radio"/>	 <a href="#">AdministratorAccess-Amplify</a>	AWS managed	0	
<input type="radio"/>	 <a href="#">AdministratorAccess-AWSElasticBeanstalk</a>	AWS managed	0	
<input type="radio"/>	 <a href="#">AIOpsAssistantPolicy</a>	AWS managed	0	
<input type="radio"/>	 <a href="#">AIOpsConsoleAdminPolicy</a>	AWS managed	0	
<input type="radio"/>	 <a href="#">AIOpsOperatorAccess</a>	AWS managed	0	

For now, we'll proceed without any permissions.

## 5. Create User

## Review and create

Review your choices. After you create the user, you can view and download the autogenerated password, if enabled.

### User details

User name  
Jeevan

Console password type  
Custom password

Require password reset  
No

### Permissions summary

< 1 >

Name 



Type



Used as



No resources

### Tags - optional

Tags are key-value pairs you can add to AWS resources to help identify, organize, or search for resources. Choose any tags you want to associate with this user.

No tags associated with the resource.

[Add new tag](#)

You can add up to 50 more tags.

[Cancel](#)

[Previous](#)

[Create user](#)

## Retrieve password

You can view and download the user's password below or email users instructions for signing in to the AWS Management Console. This is the only time you can view and download this password.

### Console sign-in details

[Email sign-in instructions !\[\]\(2885535958616e9ec6b97903614c334b\_img.jpg\)](#)


#### Console sign-in URL

 <https://537976613696.signin.aws.amazon.com/console>

#### User name

 Jeevan

#### Console password

 \*\*\*\*\* [Show](#)

[Cancel](#)

[Download .csv file](#)

[Return to users list](#)

## Console sign-in URL

<https://537976613696.signin.aws.amazon.com/console>

## User name

Jeevan

## 6. Use the link & Sign-in

<https://537976613696.signin.aws.amazon.com/console>



## IAM user sign in ⓘ

Account ID or alias [\(Don't have?\)](#)

☐ Remember this account

IAM username

Password

☐ Show Password

[Having trouble?](#)

Sign in

## Amazon Lightsail

Lightsail is the easiest way to get started on AWS

[Learn more »](#)



[Alt+S] 🔍 🔔 ⚙️ Europe (Stockholm) ▼

### Console Home ⓘ

#### Recently visited ⓘ

No recently visited services

Explore one of these commonly visited AWS services.

[EC2](#) [S3](#) [Aurora and RDS](#) [Lambda](#)

[View all services](#)

#### Applications (0) ⓘ

Region: Europe (Stockholm)

Select Region: eu-north-1 (Current Region) ▼

Name	Description	Region	Originati...	★
<div><div>🚫 Access denied to servicecatalog:ListApplications</div><div><a href="#">🔍 Diagnose with Amazon Q</a></div></div>				

[Go to myApplications](#)

#### Welcome to AWS

**Getting started with AWS** ⓘ  
Learn the fundamentals and find valuable information to get the most out of AWS.

**Training and certification** ⓘ  
Learn from AWS experts and advance your skills and knowledge.

**What's new with AWS?** ⓘ  
Discover new AWS services, features, and Regions.

#### AWS Health ⓘ

No health data

You don't have permissions to access AWS Health.

#### Cost and usage ⓘ

**Current month costs**  
🚫 Access denied

**Forecasted month end costs**  
🚫 Access denied

**Savings opportunities**  
🚫 Access denied

**Cost breakdown**  
🚫 Access denied

AWS Introduction

20



You do not have access to anything.

## 7. Give permissions

- Login your root account
- IAM → Select user

- Add permissions

Permissions Groups Tags Security credentials Last Accessed

**Permissions policies (0)**

Permissions are defined by policies attached to the user directly or through groups.

Filter by Type: All types

Search:

Policy name [↗](#) Type Attached via [↗](#)

No resources to display

[Add permissions](#) [Remove](#) [Add permissions](#) [Create inline policy](#)

▶ **Permissions boundary** (not set)

- Attach policies directly
- **Select the permissions:**  
AdministratorAccess,  
Billing,  
AWSBudgetsActions\_RolePolicyForResourceAdministrationWithSSM
- Add permissions

**Permissions summary (3)**

Name <a href="#">↗</a>	Type	Used as
<a href="#">AdministratorAccess</a>	AWS managed - job function	Permissions policy
<a href="#">Billing</a>	AWS managed - job function	Permissions policy
<a href="#">AWSBudgetsActions_RolePolicyForResourceAdministrationWithSSM</a>	AWS managed	Permissions policy

[Cancel](#) [Previous](#) [Add permissions](#)

## 8. Now login to your IAM User

EC2

Dashboard

EC2 Global View

Events

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Capacity Reservations

Images

AMIs

AMI Catalog

Elastic Block Store

Volumes

Snapshots

Resources

EC2 Global View

You are using the following Amazon EC2 resources in the Europe (Stockholm) Region:

Instances (running)	0	Auto Scaling Groups	0	Capacity Reservations	0
Dedicated Hosts	0	Elastic IPs	0	Instances	0
Key pairs	0	Load balancers	0	Placement groups	0
Security groups	1	Snapshots	0	Volumes	0

Launch instance

To get started, launch an Amazon EC2 instance, which is a virtual server in the cloud.

Launch instance

Migrate a server

Note: Your instances will launch in the Europe (Stockholm) Region

Instance alarms

View in CloudWatch

0 in alarm

0 OK

0 insufficient data

Service health

AWS Health Dashboard

Region  
Europe (Stockholm)

Status  
This service is operating normally.

Zones

Zone name	Zone ID
eu-north-1a	eun1-az1
eu-north-1b	eun1-az2



You'll see the EC2 instance



Even with **AdministratorAccess**, an IAM user does NOT have access to the **aws-portal** (billing) or other root-specific actions by default.